

Error handling design patterns in non-OOP languages

Namely in ISO C

Matjaž GUŠTIN

2019-01-09

Material

- ▶ Slides available on [**matjaz.it/slides**](https://matjaz.it/slides)
- ▶ Slides licensed under a [Creative Commons Attribution 4.0 International License \(CC BY 4.0\)](https://creativecommons.org/licenses/by/4.0/)
- ▶ LaTeX source code available at github.com/TheMatjaz/c_error_handling_design_patterns

Overview

1. A brief recap over Exceptions
2. Return codes
 - 2.1 Boolean codes
 - 2.2 Error codes
 - 2.3 Error flags
 - 2.4 Status/state codes
3. Code style for return codes
4. Using the return value's domain
5. Embedded error indicator in data type
6. `<errno.h>`
7. `<setjmp.h>`

Note: it will be very code-based and development-focused

A brief recap over Exceptions

Exceptions

In OOP languages we commonly have the **Exception** classes used to handle:

- ▶ unexpected values or states
- ▶ special cases
- ▶ non-nominal situations
- ▶ ... something that cannot be handled the normal way

Control flow break

When an Exception object is raised/thrown:

- ▶ breaks the execution flow
- ▶ makes the current function return early
- ▶ repeats the same on every function in the stack towards the `main()`
- ▶ stops if caught by a `try-catch` block
- ▶ or makes the program crash

Advantages of exceptions

- ▶ Biggest advantage: break of the control flow
- ▶ Allows separation of:
 - ▶ nominal behaviour
 - ▶ error-handling behaviour
- ▶ It's nicer to read

From *Clean Code* by Robert C. Martin

- ▶ *"Returning error codes from command functions is a subtle violation of command query separation."*
- ▶ **Command query separation:** *"functions should either do something or answer something, but not both. Either your function should change the state of an object [command], or it should return some information about that object [query]. Doing both often leads to confusion."*


```
1 def update_software():
2     completed = False
3     while not completed:
4         try:
5             update = download_latest_software_update()
6             update.check_integrity()
7             update.install()
8             completed = True
9         except ConnectionError:
10             log.error("Retrying download later.")
11             time.sleep(3600)
12         except IntegrityError:
13             log.error("Downloaded file corrupted.")
14         except InstallationError:
15             log.error("No admin rights.")
16             print("Please run the program with sudo.")
17             completed = True
18         finally:
19             update.erase_temp_files()
```

But...

- ▶ Many programming languages were created before OOP was even a thing (duh...)
- ▶ Examples: Fortran, C, Cobol, ... or the more modern Rust
- ▶ Other patterns took place, still in use today
- ▶ Main difference: **the control flow is not broken!** Manual handling is required.
- ▶ *"If a tree falls in a forest and no one is around to hear it, does it make a sound?"*

Problem and forces

- ▶ Need a way to indicate to the function caller that something could not be done and why.
- ▶ Exceptions are not available.
- ▶ Must be simple, lightweight, efficient, easy to understand.

Overview of the patterns

1. Return codes

1.1 Boolean codes

1.2 Error codes

1.3 Error flags

1.4 Status/state codes

2. Using the return value's domain

3. Embedded error indicator in data type

4. `<errno.h>`

5. `<setjmp.h>`

Return codes

Return codes

- ▶ The function's return value indicates its execution success or failure
- ▶ Different levels of detail in the output
- ▶ A potential way to categorize them¹
 1. Boolean codes
 2. Error codes
 3. Error flags
 4. Status/state codes

¹This is just my proposal, as often all these terms are used interchangeably

Boolean return code

```
1 #include <stdbool.h>
2
3 bool receive_message(message_t* message);
4
5 // Alternate version without booleans
6 int receive_message(message_t* message);
```

true (or non-0) on success

false (or 0) on failure

Boolean return code: consequences

- Why did it fail?
- Can we retry or not? Maybe with different settings?
- Confusing: does `false` indicate "no error" or "no success"?

Boolean return code: usage examples

```
1 message_t rx_message;  
2 if (!receive_message(&rx_message)) {  
3     puts("Reception failure");  
4 }
```

```
1 message_t rx_message;  
2 while (!receive_message(&rx_message)) {  
3     sleep(5);  
4 }
```

Hard to remember if the negation is needed or not.

Error code

```
1 typedef enum {  
2     RX_OK = 0,  
3     ERROR_TIMEOUT_NOTHING_RECEIVED,  
4     ERROR_BROKEN_CRC,  
5     ERROR_INCOMPLETE_MESSAGE,  
6     ERROR_ANTENNA_DISCONNECTED,  
7 } rx_code_t;  
8  
9 rx_code_t receive_message(message_t* message);
```

RX_OK on success. OK is false to indicate the *absence* of errors.

Others on failure. The actual value indicates the exact reason.

Error code: consequences

- Cleaner setup
- We can handle different failure cases in different ways
- Easy to remember which value is the success: the `false` value
- Can also be used in binary/encoded/serialized messages

```
1 typedef struct {  
2     uint32_t error_code; // Indicate if the value makes sense  
3     uint32_t value;      // Value to transmit  
4 } packed_value_t;
```

- **Arguably the most common pattern** outside of `libc`
- Longer code handling the cases

Usage examples: error code

```
1 message_t rx_message;
2 rx_code_t error;
3
4 error = receive_message(&rx_message);
5 if (error != RX_OK) { // Simply: if (error) {...}
6     puts("Reception failure");
7 }
```

```
1 message_t rx_message;
2 rx_code_t rx_code;
3
4 if ((rx_code = receive_message(&rx_message)) != RX_OK) {
5     puts("Reception failure");
6 }
```

```

1 message_t rx_message;
2 rx_code_t rx_code;
3 bool keep_receiving = true;
4
5 do {
6     rx_code = receive_message(&rx_message);
7     switch (rx_code) {
8         case (RX_OK): {
9             process_message(&rx_message);
10            break;
11        }
12        case (ERROR_TIMEOUT_NOTHING_RECEIVED) {
13            sleep(5);
14            break;
15        }
16        case (ERROR_ANTENNA_DISCONNECTED) {
17            puts("Please connect the antenna to the PC");
18            keep_receiving = false;
19            break;
20        }
21        default: {} // Do nothing, just retry the reception.
22    }
23 } while (keep_receiving);

```

Usage examples: process exit status

- ▶ Also known as *exit code* or *error level*.
- ▶ Value that a process returns after its termination to the parent process.
- ▶ Usually `int32` where 0 means OK: process completed successfully.
- ▶ Non-zero values are not standardized: every OS has a different list of recommended/preferred interpretations.
- ▶ In C: the `int` value returned by `main()`. Alternatively the argument of `exit()`.

Error flags

```
1 typedef enum {  
2     // Bit flags, each value on different bit  
3     RX_OK = 0x00,  
4     ERROR_TIMEOUT_NOTHING_RECEIVED = 0x01,  
5     ERROR_BROKEN_CRC = 0x02,  
6     ERROR_INCOMPLETE_MESSAGE = 0x04,  
7     ERROR_ANTENNA_DISCONNECTED = 0x08,  
8 } rx_flag_t;  
9 typedef uint8_t rx_code_t;  
10  
11 rx_code_t receive_message(message_t* message);
```

RX_OK on success. No flags or false to indicate the *absence* of errors.

Any flag on failure. Each bit expresses one reason. More than one reason possible **simultaneously**.

Error flags: consequences

In addition to the consequences of Error codes:

- Useful if multiple failures can happen simultaneously
- Can still be used if only one failure at the time
- Easy to remember which value is the success: the false value
- N bits indicate only N errors. Error codes indicate $2^N - 1$. Bigger integer types may be needed.
- Even longer code handling the cases (e.g. a series of `if-if-if` but not `switch-case`)
- Macros are required to group same-category flags

Status code

```
1 typedef enum {
2     STATUS_OK_DETERMINISTIC = 0,
3     STATUS_OK_PROBABILISTIC_HIGH_PROB,
4     STATUS_OK_PROBABILISTIC_LOW_PROB,
5     STATUS_ERROR_ILLEGAL_HEADER,
6     STATUS_ERROR_ILLEGAL_CONTENT,
7     STATUS_ERROR_ILLEGAL_ENCODING,
8 } status_code_t;
9 #define STATUS_IS_OK(s) ((s) == STATUS_OK_DETERMINISTIC || (s) ==
10     STATUS_OK_PROBABILISTIC_HIGH_PROB || (s) ==
11     STATUS_OK_PROBABILISTIC_LOW_PROB )
12
13 status_code_t process_message(message_t* message);
```

Multiple ways to complete a task (e.g. Strategy) and we know which one was chosen.

Status code: consequences

- We can handle different success cases in different ways
- Easily confusing
- Often macros are used to group "OK"-states, which results in a regular error code

Status code split setup

```
1 typedef enum {
2     STATUS_OK = 0,
3     STATUS_ERROR_ILLEGAL_HEADER,
4     STATUS_ERROR_ILLEGAL_CONTENT,
5     STATUS_ERROR_ILLEGAL_ENCODING,
6 } error_code_t;
7 typedef enum {
8     STATUS_EXTRA_DETERMINISTIC = 0,
9     STATUS_EXTRA_PROBABILISTIC_HIGH_PROB,
10    STATUS_EXTRA_PROBABILISTIC_LOW_PROB,
11 } status_extra_t;
12
13 error_code_t process_message(message_t* message,
14                             status_extra_t* extra);
```

Status code split setup: consequences

- We separated the execution success/failure indicator from the additional information
- Now we need to handle two values instead of one
- What is the value of `extra` if an error happens during processing?
- What happens if I pass a NULL pointer for `extra`?

Nested return codes

```
1 rx_code_t receive_message(message_t* message) {  
2     rx_code_t reception_error = RX_OK;  
3     encoding_code_t encoding_error = ENC_OK;  
4  
5     encoding_error = prepare_message(message);  
6     if (encoding_error) {  
7         return ???; // Which error code should we return?  
8     }  
9     ...  
10 }
```

Nested return codes problem (cont.)

- ▶ Returning the inner error code `encoding_error` breaks abstraction layers
- ▶ Returning the outer error code `reception_error` may hide details
- ▶ Using only one enum for everything: very big enum?
- ▶ Combining the error code into a concatenated integer leads to confusion

Nested return codes problem (cont.)

Common solutions: break abstraction layers

- ▶ Either one huge enum containing every possible error
 - ▶ Each function inside the library uses the same error code data type
 - ▶ Just returning it to the upper layer
- ▶ Or combined error code
 - ▶ Example: ISO/IEC 7816 for smart cards uses `uint16`
 - ▶ High byte for category (where failed)
 - ▶ Low byte for reason (what failed)
 - ▶ `0x9000` = OK
 - ▶ `0x69xx` = Something with the command
 - ▶ `0x6900` = Command not allowed
 - ▶ `0x6981` = Command not compatible with data structure

Code style for return codes

Stopping flow with multiple return

```
1 rx_code_t receive_message(message_t* message) {
2     receiver_config_t config;  rx_code_t error_code;
3
4     error_code = load_receiver_config_from_storage(&config);
5     if (error_code) return error_code;  // Bad, missing brackets
6     error_code = receiver_enable_peripheral();
7     if (error_code) { return error_code; }  // Always use brackets
8     error_code = receiver_configure(&config);
9     if (error_code) { return error_code; }
10    error_code = receiver_receive(message);
11    return error_code;
12 }
```

With `do-while(0)` is nicer

```
1 rx_code_t receive_message(message_t* message) {  
2     receiver_config_t config;  rx_code_t error_code;  
3  
4     do {  
5         error_code = load_receiver_config_from_storage(&config);  
6         if (error_code) { break; }  
7         error_code = receiver_enable_peripheral();  
8         if (error_code) { break; }  
9         error_code = receiver_configure(&config);  
10        if (error_code) { break; }  
11        error_code = receiver_receive(message);  
12    } while (0);  
13    return error_code;  
14 }
```

Loops inside do-while(0) don't work

```
1 rx_code_t receive_message(message_t* message) {
2     receiver_config_t config;  rx_code_t error_code;
3     do {
4         error_code = load_receiver_config_from_storage(&config);
5         if (error_code) { break; }
6         for (int i = 0; i < 10; i++) {
7             // Activate 10 peripherals
8             error_code = receiver_enable_peripheral(i);
9             if (error_code) { break; }
10            // The break exits the for loop, not the while(0)
11        }
12        error_code = receiver_configure(&config);
13        if (error_code) { break; }
14        error_code = receiver_receive(message);
15    } while (0);
16    return error_code;
17 }
```

With goto works always

```
1 rx_code_t receive_message(message_t* message) {
2     receiver_config_t config; rx_code_t error_code;
3     error_code = load_receiver_config_from_storage(&config);
4     if (error_code) { goto terminate; }
5     for (int i = 0; i < 10; i++) {
6         // Activate 10 peripherals
7         error_code = receiver_enable_peripheral(i);
8         if (error_code) { goto terminate; }
9     }
10    error_code = receiver_configure(&config);
11    if (error_code) { goto terminate; }
12    error_code = receiver_receive(message);
13
14    terminate: {
15        return error_code;
16    }
17 }
```

My personal rules for goto

1. Avoid it, if you can.
2. Use only **one label per function** with a clear name:
 - ▶ Good: `goto termination`, `goto error_handling`
 - ▶ Bad: `goto failure` (Does it jump to a failing point or handles a failure?)
3. Only jump **downwards in the code** (i.e. skip some instruction, don't cycle them). Better if only to the function's end (to the return call).

Comparing various code styles

- ▶ Multiple `return`: clear when writing, confusing to read. Not allowed by certain industry standards (e.g. automotive).
- ▶ `do-while(0)` is clear for simple cases after getting used to it
- ▶ `goto` is clear only when used correctly and only for this kind of handling

Using the return value's domain

Return values outside the domain

- ▶ The function returns a value, not a return code.
- ▶ The value has a limited domain.
- ▶ When value out of bounds, indicates an error.

Example: writing formatted strings to a file.

```
1 int fprintf ( FILE * stream, const char * format, ... );
```

Returns the amount of characters written: 0 or more.

Negative on failure.

Outside the domain: consequences

- No need for additional enums
- Easy to understand if something is wrong (e.g. negative length does not make sense)
- Easy to forget to check and use error value as a good result
- Must read documentation of function in detail
- Not possible if no value outside the domain exists

Return values inside the domain

- ▶ Also known as *semipredicate problem*
- ▶ how do I distinguish between a valid output and an error indicator?

Example: parsing a string for an integer value.

```
1 long int strtol (const char* str, char** endptr, int base);
```

Returns the converted integer.

If no valid conversion could be performed, **returns 0**.

Return values inside the domain (cont.)

What happens if the input string contains a zero digit?

```
1 strtol("12", NULL, 10); // returns 12
2
3 strtol("0", NULL, 10);  // returns 0
4
5 strtol("fdjkfnskxg", NULL, 10); // returns 0 as well!
```

But: `strtol()` sets second parameter `endptr` to first char after the parsed number. On failed parsing `str == endptr`.

Inside the domain: consequences

- Extremely confusing
- Just never ever do that

Embedded error indicator in data type

Nullable types

The language's type system supports every value to be either NULL-like indicating missing data or a value.

- ▶ In Python anything can be None
- ▶ In SQL anything can be NULL
- ▶ In Java non-primitives only (int no, Integer yes): *Null object pattern*
- ▶ In C works only with pointers

Nullable types: example from SQL

ISO SQL (structured query language) used to interrogate RDBMS:

- ▶ every value can also be NULL indicating missing data
- ▶ functions often return either the value or NULL
- ▶ three-valued logic with NULL as "undefined":
 - ▶ NULL AND TRUE → NULL
 - ▶ NULL AND FALSE → FALSE
 - ▶ NULL AND NULL → NULL
 - ▶ NULL OR NULL → NULL

A simple pointer

A pointer to something may be NULL (have the value 0) to indicate the broken link. Otherwise it can be dereferenced.

```
1 #include <stdio.h>
2 #include <stdint.h>
3
4 int main() {
5     uint8_t* buffer = malloc(50);
6     if (buffer != NULL) {
7         puts("Malloc succeeded and I have a buffer");
8     } else {
9         puts("Malloc failed");
10    }
11    return 0;
12 }
```


A simple pointer: consequences

- Embedded in most programming languages
- Every programmer understands: cannot work on NULL
- No need for checking a separate error code
- Very easy to forget to check and dereference NULL (auch!)
- Dereferencing may be too slow
- In embedded environments may be possible work due to missing `malloc()`
- May make the code less readable (`int*` vs `int`)

Option(al) types

Polymorphic type representing a value that may or may not have meaning. Mostly functional languages.

- ▶ In Rust: `enum Option<T> { None, Some(T) }`
- ▶ In Haskell: `data Maybe a = Nothing | Just a`
- ▶ In C#: `Nullable<T>` or `T?`
- ▶ In C: manually...

```
1 struct optional_uint32 {  
2     uint8_t error_code; // Indicate if the value makes sense  
3     uint32_t value;      // Actual value  
4 };
```

IEEE 754 floating-point values

The IEEE encoding of the real numbers has embedded handling of special cases.

▶ $+\infty$

- ▶ division by (positive) zero: `1.0/0.0`
- ▶ overflows: `pow(10.0, 500.0)`
- ▶ always bigger than any other value
- ▶ to check: `isinf(value) && value > 0`

▶ $-\infty$

- ▶ division of negative value by zero: `-1.0/0.0`
- ▶ or by negative zero: `1.0/-0.0`
- ▶ underflows: `pow(-10.0, 501.0)`
- ▶ always smaller than any other value
- ▶ to check: `isinf(value) && value < 0`

IEEE 754 floating-point values (cont.)

- ▶ NaN (Not a Number)
 - ▶ invalid operations
 - ▶ $0.0/0.0$
 - ▶ $\text{Infinity} * 0.0$
 - ▶ $\text{sqrt}(-1.0)$
 - ▶ propagating: operations on a NaN return NaN
 - ▶ comparing with a NaN always returns false
 - ▶ has a quiet (just returning) and signaling variant (FPU/floatlib indicates error to the system)
 - ▶ it even has free bits to encode type of error (rare)
 - ▶ to check: `isnan(value)`
- ▶ combined check: `isfinite(value)` returns true when the value is not NaN or $\pm\infty$

IEEE 754 floating-point values: consequences

- IEEE standard since decades, really every computer supports it
- Hardware accelerated
- No need for pointer dereferencing or checking a separate error code
- Easy to forget to check with `isfinite()`

<errno.h>

<errno.h>

A standard, glorified global error code.

```
1 #include <stdio.h>
2 #include <errno.h>    // To access the integer `errno`
3 #include <string.h>    // To convert `errno` to a human-readable
   string with strerror()
4
5 int main() {
6     printf("At startup: value=%d, string=%s\n",
7         errno, strerror(errno));
8     FILE *file = fopen("NON_existing_file.txt", "r");
9     if (file == NULL) {
10         printf("After fopen fails: value=%d, string=%s\n",
11             errno, strerror(errno));
12     } else { fclose(file); }
13     return 0;
14 }
```

<errno.h>: consequences

- Part of C standard library
- Easy to lose track who set the `errno` variable in nested code
- Global variables are bad

<setjmp.h>

<setjmp.h>

```
1 int setjmp (jmp_buf env);
```

- ▶ Fills env with the current state of the calling environment, so it can be restored later.
- ▶ Returns 0 on direct invocation (when state is saved).
- ▶ Otherwise returns the value (forcibly non-zero) passed by longjmp() (when state is restored).

<setjmp.h> (cont.)

```
1 void longjmp (jmp_buf env, int val);
```

- ▶ Restores stored env.
- ▶ Transfers the control to the point where setjmp() was last used to fill the env.
- ▶ Makes setjmp() return val.
- ▶ This function never returns (jumping to setjmp() before that).

```
1 #include <stdio.h>
2 #include <setjmp.h>
3
4 static jmp_buf state;
5 typedef enum {
6     OK = 0,
7     NEGATIVE_VALUE = 1,
8     TOO_BIG_VALUE = 2,
9 } error_code_t;
10
11 int twice4(int value) {
12     if (value < 0) {
13         printf("(!) Negative value: %d\n", value);
14         longjmp(state, NEGATIVE_VALUE); // Restore state, set code
15     } else if (value > 100) {
16         printf("(!) Too big value: %d\n", value);
17         longjmp(state, TOO_BIG_VALUE); // Restore state, set code
18     } else { return 2 * value; }
19 }
```

```
1 int twice3(int value) { return twice4(value); }
2 int twice2(int value) { return twice3(value); }
3 int twice(int value) { return twice2(value); }
4
5 int main() {
6     // Initially saves state and sets error_code to 0.
7     // Jumped to using longjmp(state, new_value),
8     // setting error_code to new_value.
9     int error_code = setjmp(state);
10    if (error_code == 0K) {
11        int input = -10; // -10 or 1000 jumps to else branch
12        int result = twice(input);
13        printf("Twice of %d is %d\n", input, result);
14    } else {
15        printf("Error code %d\n", error_code);
16    }
17    return 0;
18 }
```

<setjmp.h>: consequences

- Breaking control flow
- Good performance (avoiding functions return calls)
- A way to implement exception-like behaviour
- Like goto but worse: may be **very** confusing
- Often readability is more important than premature optimization

Conclusion

Wrapping up

- ▶ If your programming language supports nullable/option types: use them
- ▶ Otherwise go with return error codes or flags: 0 for OK, other values for error cases
- ▶ Write the error handling code with care, focus on **readability**

Sources

- ▶ <http://www.cplusplus.com/reference/>
- ▶ https://en.wikipedia.org/wiki/Exception_handling
- ▶ https://en.wikipedia.org/wiki/Exit_status
- ▶ https://en.wikipedia.org/wiki/Nullable_type
- ▶ https://en.wikipedia.org/wiki/Floating-point_arithmetic#Special_values
- ▶ Robert C. Martin, *Clean Code: A Handbook of Agile Software Craftsmanship*, ISBN-10: 9780132350884

Material

- ▶ Slides available on [**matjaz.it/slides**](https://matjaz.it/slides)
- ▶ Slides licensed under a [Creative Commons Attribution 4.0 International License \(CC BY 4.0\)](https://creativecommons.org/licenses/by/4.0/)
- ▶ LaTeX source code available at github.com/TheMatjaz/c_error_handling_design_patterns