

Error handling design patterns in non-OOP languages

Namely in ISO C

Matjaž GUŠTIN

2019-01-09

Material

- ▶ Slides available on [**matjaz.it/slides**](https://matjaz.it/slides)
- ▶ Slides licensed under a [Creative Commons Attribution 4.0 International License \(CC BY 4.0\)](https://creativecommons.org/licenses/by/4.0/)
- ▶ LaTeX source code available at github.com/TheMatjaz/c_error_handling_design_patterns

Overview

1. A brief recap over Exceptions
2. Return codes
 - 2.1 Boolean codes
 - 2.2 Error codes
 - 2.3 Error flags
3. Using the return value's domain
4. Embedded error indicator in data type
5. `<errno.h>`
6. `<setjmp.h>`
7. Code design choices with return codes

Note: it will be very code-based and development-focused

A brief recap over Exceptions

Exceptions

In OOP languages we commonly have the **Exception** classes used to handle:

- ▶ unexpected values or states
- ▶ special cases
- ▶ non-nominal situations
- ▶ ... something that cannot be handled the normal way

Problem and forces

- ▶ Need a way to indicate to the function caller that something could not be done and why.
- ▶ Exceptions are not available.
- ▶ Must be simple, lightweight, efficient, easy to understand.

Return codes

Return codes

- ▶ Also known as *status codes*
- ▶ The function's return value indicates its execution success or failure
- ▶ Different levels of detail
- ▶ A potential way to categorize them¹
 1. Boolean codes
 2. Error codes
 3. Error flags

¹This is just my proposal, as often all these terms are used interchangeably

Boolean return code

```
1 #include <stdbool.h>
2
3 bool receive_message(message_t* message);
4
5 // Alternate version without booleans
6 int receive_message(message_t* message);
```

true (or non-0) on success

false (or 0) on failure

Boolean return code: consequences

- Confusing: does `false` indicate "no error" or "no success"?
- Why did it fail?
- Can we retry or not? Maybe with different settings?

Error code

```
1 typedef enum {  
2     RX_OK = 0,  
3     ERROR_TIMEOUT_NOTHING_RECEIVED,  
4     ERROR_BROKEN_CRC,  
5     ERROR_INCOMPLETE_MESSAGE,  
6     ERROR_ANTENNA_DISCONNECTED,  
7 } rx_code_t;  
8  
9 rx_code_t receive_message(message_t* message);
```

RX_OK on success. OK is false to indicate the *absence* of errors.

Others on failure. The actual value indicates the exact reason.

Error code: consequences

- Cleaner setup
- We can handle different failure cases in different ways
- Easy to remember which value is the success: the `false` value
- **Arguably the most common pattern** outside of `libc`
- Longer code handling the cases

Usage examples: error code

```
1 message_t rx_message;  
2 rx_code_t error;  
3  
4 error = receive_message(&rx_message);  
5 if (error != RX_OK) { // Simply: if (error) {...}  
6     puts("Reception failure");  
7 }
```

```

1 message_t rx_message;
2 rx_code_t rx_code;
3 bool keep_receiving = true;
4
5 do {
6     rx_code = receive_message(&rx_message);
7     switch (rx_code) {
8         case (RX_OK): {
9             process_message(&rx_message);
10            break;
11        }
12        case (ERROR_TIMEOUT_NOTHING_RECEIVED) {
13            sleep(5);
14            break;
15        }
16        case (ERROR_ANTENNA_DISCONNECTED) {
17            puts("Please connect the antenna to the PC");
18            keep_receiving = false;
19            break;
20        }
21        default: {} // Do nothing, just retry the reception.
22    }
23 } while (keep_receiving);

```

Usage examples: process exit status

- ▶ Also known as *exit code* or *error level*.
- ▶ Value that a process returns after its termination to the parent process.
- ▶ Usually `int32` where 0 means OK: process completed successfully.
- ▶ Non-zero values are not standardized: every OS has a different list of recommended/preferred interpretations.
- ▶ In C: the `int` value returned by `main()`. Alternatively the argument of `exit()`.

Error flags

```
1 typedef enum {  
2     // Bit flags, each value on different bit  
3     RX_OK = 0x00,  
4     ERROR_TIMEOUT_NOTHING_RECEIVED = 0x01,  
5     ERROR_BROKEN_CRC = 0x02,  
6     ERROR_INCOMPLETE_MESSAGE = 0x04,  
7     ERROR_ANTENNA_DISCONNECTED = 0x08,  
8 } rx_flag_t;  
9 typedef uint8_t rx_code_t;  
10  
11 rx_code_t receive_message(message_t* message);
```

RX_OK on success. No flags or false to indicate the *absence* of errors.

Any flag on failure. Each bit expresses one reason. More than one reason possible **simultaneously**.

Error flags: consequences

In addition to the consequences of Error codes:

- Useful if multiple failures can happen simultaneously
- N bits indicate only N errors. Error codes indicate $2^N - 1$. Bigger integer types may be needed.
- Even longer code handling the cases: need to handle all possible flags independently (e.g. a series of `if-if-if` but not `switch-case`)
- Sometimes macros are involved for operations on groups of flags

Using the return value's domain

Return values outside the domain

- ▶ The function returns a value, not a return code.
- ▶ The value has a limited domain.
- ▶ When value out of bounds, indicates an error.

Example: writing formatted strings to a file.

```
1 int fprintf ( FILE * stream, const char * format, ... );
```

Returns the amount of characters written: 0 or more.

Negative on failure.

Outside the domain: consequences

- No need for additional enums
- Easy to understand if something is wrong (e.g. negative length does not make sense)
- Easy to forget to check and use error value as a good result
- Must read documentation of function in detail
- Not possible if no value outside the domain exists

Embedded error indicator in data type

Nullable types

The language's type system supports every value to be either NULL-like indicating missing data or a value.

- ▶ In Python anything can be None
- ▶ In SQL anything can be NULL
- ▶ In Java non-primitives only (int no, Integer yes): *Null object pattern*
- ▶ In C works only with pointers

A simple pointer

A pointer to something may be NULL (have the value 0) to indicate the broken link. Otherwise it can be dereferenced.

```
1 #include <stdio.h>
2 #include <stdint.h>
3
4 int main() {
5     uint8_t* buffer = malloc(50);
6     if (buffer != NULL) {
7         puts("Malloc succeeded and I have a buffer");
8     } else {
9         puts("Malloc failed");
10    }
11    return 0;
12 }
```

A simple pointer: consequences

- Embedded in most programming languages
- Every programmer understands: cannot work on NULL
- No need for checking a separate error code
- Very easy to forget to check and dereference NULL (auch!)
- Dereferencing may be too slow
- In embedded environments may be possible work due to missing `malloc()`
- May make the code less readable (`int*` vs `int`)

Option(al) types

Polymorphic type representing a value that may or may not have meaning. Mostly functional languages.

- ▶ In Rust: `enum Option<T> { None, Some(T) }`
- ▶ In Haskell: `data Maybe a = Nothing | Just a`
- ▶ In C#: `Nullable<T>` or `T?`
- ▶ In C: manually...

```
1 struct optional_uint32 {  
2     uint8_t error_code; // Indicate if the value makes sense  
3     uint32_t value;      // Actual value  
4 };
```

IEEE 754 floating-point values

The IEEE encoding of the real numbers has embedded handling of special cases.

▶ $+\infty$

- ▶ division by (positive) zero: `1.0/0.0`
- ▶ overflows: `pow(10.0, 500.0)`
- ▶ always bigger than any other value
- ▶ to check: `isinf(value) && value > 0`

▶ $-\infty$

- ▶ division of negative value by zero: `-1.0/0.0`
- ▶ or by negative zero: `1.0/-0.0`
- ▶ underflows: `pow(-10.0, 501.0)`
- ▶ always smaller than any other value
- ▶ to check: `isinf(value) && value < 0`

IEEE 754 floating-point values (cont.)

- ▶ NaN (Not a Number)
 - ▶ invalid operations
 - ▶ $0.0/0.0$
 - ▶ $\text{Infinity} * 0.0$
 - ▶ $\text{sqrt}(-1.0)$
 - ▶ propagating: operations on a NaN return NaN
 - ▶ comparing with a NaN always returns false
 - ▶ has a quiet (just returning) and signaling variant (FPU/floatlib indicates error to the system)
 - ▶ it even has free bits to encode type of error (rare)
 - ▶ to check: `isnan(value)`
- ▶ combined check: `isfinite(value)` returns true when the value is not NaN or $\pm\infty$

IEEE 754 floating-point values: consequences

- IEEE standard since decades, really every computer supports it
- Hardware accelerated
- No need for pointer dereferencing or checking a separate error code
- Easy to forget to check with `isfinite()`

<errno.h>

<errno.h>

A standard, glorified global error code.

```
1 #include <stdio.h>
2 #include <errno.h>    // To access the integer `errno`
3 #include <string.h>    // To convert `errno` to a human-readable
   string with strerror()
4
5 int main() {
6     printf("At startup: value=%d, string=%s\n",
7           errno, strerror(errno));
8     FILE *file = fopen("NON_existing_file.txt", "r");
9     if (file == NULL) {
10         printf("After fopen fails: value=%d, string=%s\n",
11               errno, strerror(errno));
12     } else { fclose(file); }
13     return 0;
14 }
```

<errno.h>: consequences

- Part of C standard library
- Easy to lose track who set the `errno` variable in nested code
- Global variables are bad

<setjmp.h>

<setjmp.h>

```
1 int setjmp (jmp_buf env);
```

- ▶ Fills env with the current state of the calling environment, so it can be restored later.
- ▶ Returns 0 on direct invocation (when state is saved).
- ▶ Otherwise returns the value (forcibly non-zero) passed by longjmp() (when state is restored).

<setjmp.h> (cont.)

```
1 void longjmp (jmp_buf env, int val);
```

- ▶ Restores stored env.
- ▶ Transfers the control to the point where `setjmp()` was last used to fill the env.
- ▶ Makes `setjmp()` return `val`.
- ▶ This function never returns (jumping to `setjmp()` before that).

```
1 #include <stdio.h>
2 #include <setjmp.h>
3
4 static jmp_buf state;
5 typedef enum {
6     OK = 0,
7     NEGATIVE_VALUE = 1,
8     TOO_BIG_VALUE = 2,
9 } error_code_t;
10
11 int twice4(int value) {
12     if (value < 0) {
13         printf("(!) Negative value: %d\n", value);
14         longjmp(state, NEGATIVE_VALUE); // Restore state, set code
15     } else if (value > 100) {
16         printf("(!) Too big value: %d\n", value);
17         longjmp(state, TOO_BIG_VALUE); // Restore state, set code
18     } else { return 2 * value; }
19 }
```

```
1 int twice3(int value) { return twice4(value); }
2 int twice2(int value) { return twice3(value); }
3 int twice(int value) { return twice2(value); }
4
5 int main() {
6     // Initially saves state and sets error_code to 0.
7     // Jumped to using longjmp(state, new_value),
8     // setting error_code to new_value.
9     int error_code = setjmp(state);
10    if (error_code == 0K) {
11        int input = -10; // -10 or 1000 jumps to else branch
12        int result = twice(input);
13        printf("Twice of %d is %d\n", input, result);
14    } else {
15        printf("Error code %d\n", error_code);
16    }
17    return 0;
18 }
```

<setjmp.h>: consequences

- Breaking control flow
- Good performance (avoiding functions return calls)
- A way to implement exception-like behaviour
- Like goto but worse: may be **very** confusing
- Often readability is more important than premature optimization

Code design choices with return codes

Nested return codes problem

```
1 tx_code_t transmit_message(message_t* message) {  
2     tx_code_t tx_error = TX_OK;  
3     encoding_code_t encoding_error = ENC_OK;  
4  
5     encoding_error = prepare_message(message);  
6     if (encoding_error) {  
7         return ???; // Which error code should we return?  
8     }  
9     ...  
10 }
```

Nested return codes problem (cont.)

- ▶ Returning the inner error code `encoding_error` breaks abstraction layers
- ▶ Returning the outer error code `tx_error` may hide details

Nested return codes: solution 1

One huge enum containing every possible error
(Example: SQLite < v3.3.8)

```
1 typedef enum {
2     TX_OK = 0,
3     ERROR_ENCODING_HEADER,
4     ERROR_ENCODING_BODY,
5     ERROR_WRONG_CONFIGURATION,
6     ERROR_ANTENNA_DISCONNECTED,
7 } tx_code_t;
8
9 tx_code_t transmit_message(message_t* message) {
10     tx_code_t tx_error = RX_OK;
11
12     tx_error = prepare_message(message);
13     if (tx_error) { return tx_error; }
14     ...
15 }
```

Solution 1: consequences

- Every library function (inner and outer) returns the same data type
- Easy to write library: on error, just pass error code to caller
- Easy to write application: only one enum to handle
- Abstraction layers are broken
- Hard to understand which function from an API may return which codes

Nested return codes: solution 2

Combined error codes

(Example: ISO/IEC 7816 for smart cards)

```
1 typedef enum {  
2     TX_OK = 0,  
3     ERR_ENCODING,  
4     ERR_ANTENNA,  
5 } tx_categ_t;
```

```
1 typedef enum {  
2     ENC_OK = 0,  
3     ENC_HEADER,  
4     ENC_BODY,  
5 } tx_encoding_t;
```

```
1 typedef enum {  
2     ANT_OK = 0,  
3     ANT_DISCONN,  
4 } tx_antenna_t;
```

```
1 uint16_t transmit_message(message_t* message) {  
2     uint8_t tx_error_low = 0;  
3     tx_error_low = prepare_message(message);  
4     if (tx_error_low) { return (ERR_ENCODING<<8) | tx_error_low; }  
5     ...
```

Solution 2: consequences

- Abstraction layers are "less" broken: every byte is on its own layer
- Easy to understand which function from an API may return which codes
- Error codes need to be combined before returning
- Error codes need to be unpacked before inspection
- The application may ignore the "more detailed byte"

Conclusion

Wrapping up

- ▶ If your programming language supports exceptions or nullable/option types: use them
- ▶ Otherwise go with return error codes or flags: 0 for OK, other values for error cases
- ▶ Write the error handling code with care, focus on **readability**

Sources

- ▶ <http://www.cplusplus.com/reference/>
- ▶ https://en.wikipedia.org/wiki/Exception_handling
- ▶ https://en.wikipedia.org/wiki/Exit_status
- ▶ https://en.wikipedia.org/wiki/Nullable_type
- ▶ https://en.wikipedia.org/wiki/Floating-point_arithmetic#Special_values
- ▶ Robert C. Martin, *Clean Code: A Handbook of Agile Software Craftsmanship*, ISBN-10: 9780132350884

Material

- ▶ Slides available on [**matjaz.it/slides**](https://matjaz.it/slides)
- ▶ Slides licensed under a [Creative Commons Attribution 4.0 International License \(CC BY 4.0\)](https://creativecommons.org/licenses/by/4.0/)
- ▶ LaTeX source code available at github.com/TheMatjaz/c_error_handling_design_patterns